


Performance Evaluation of LLM Inference Accelerators



(a) Inference Latency Comparison



(b) Energy Consumption Breakdown